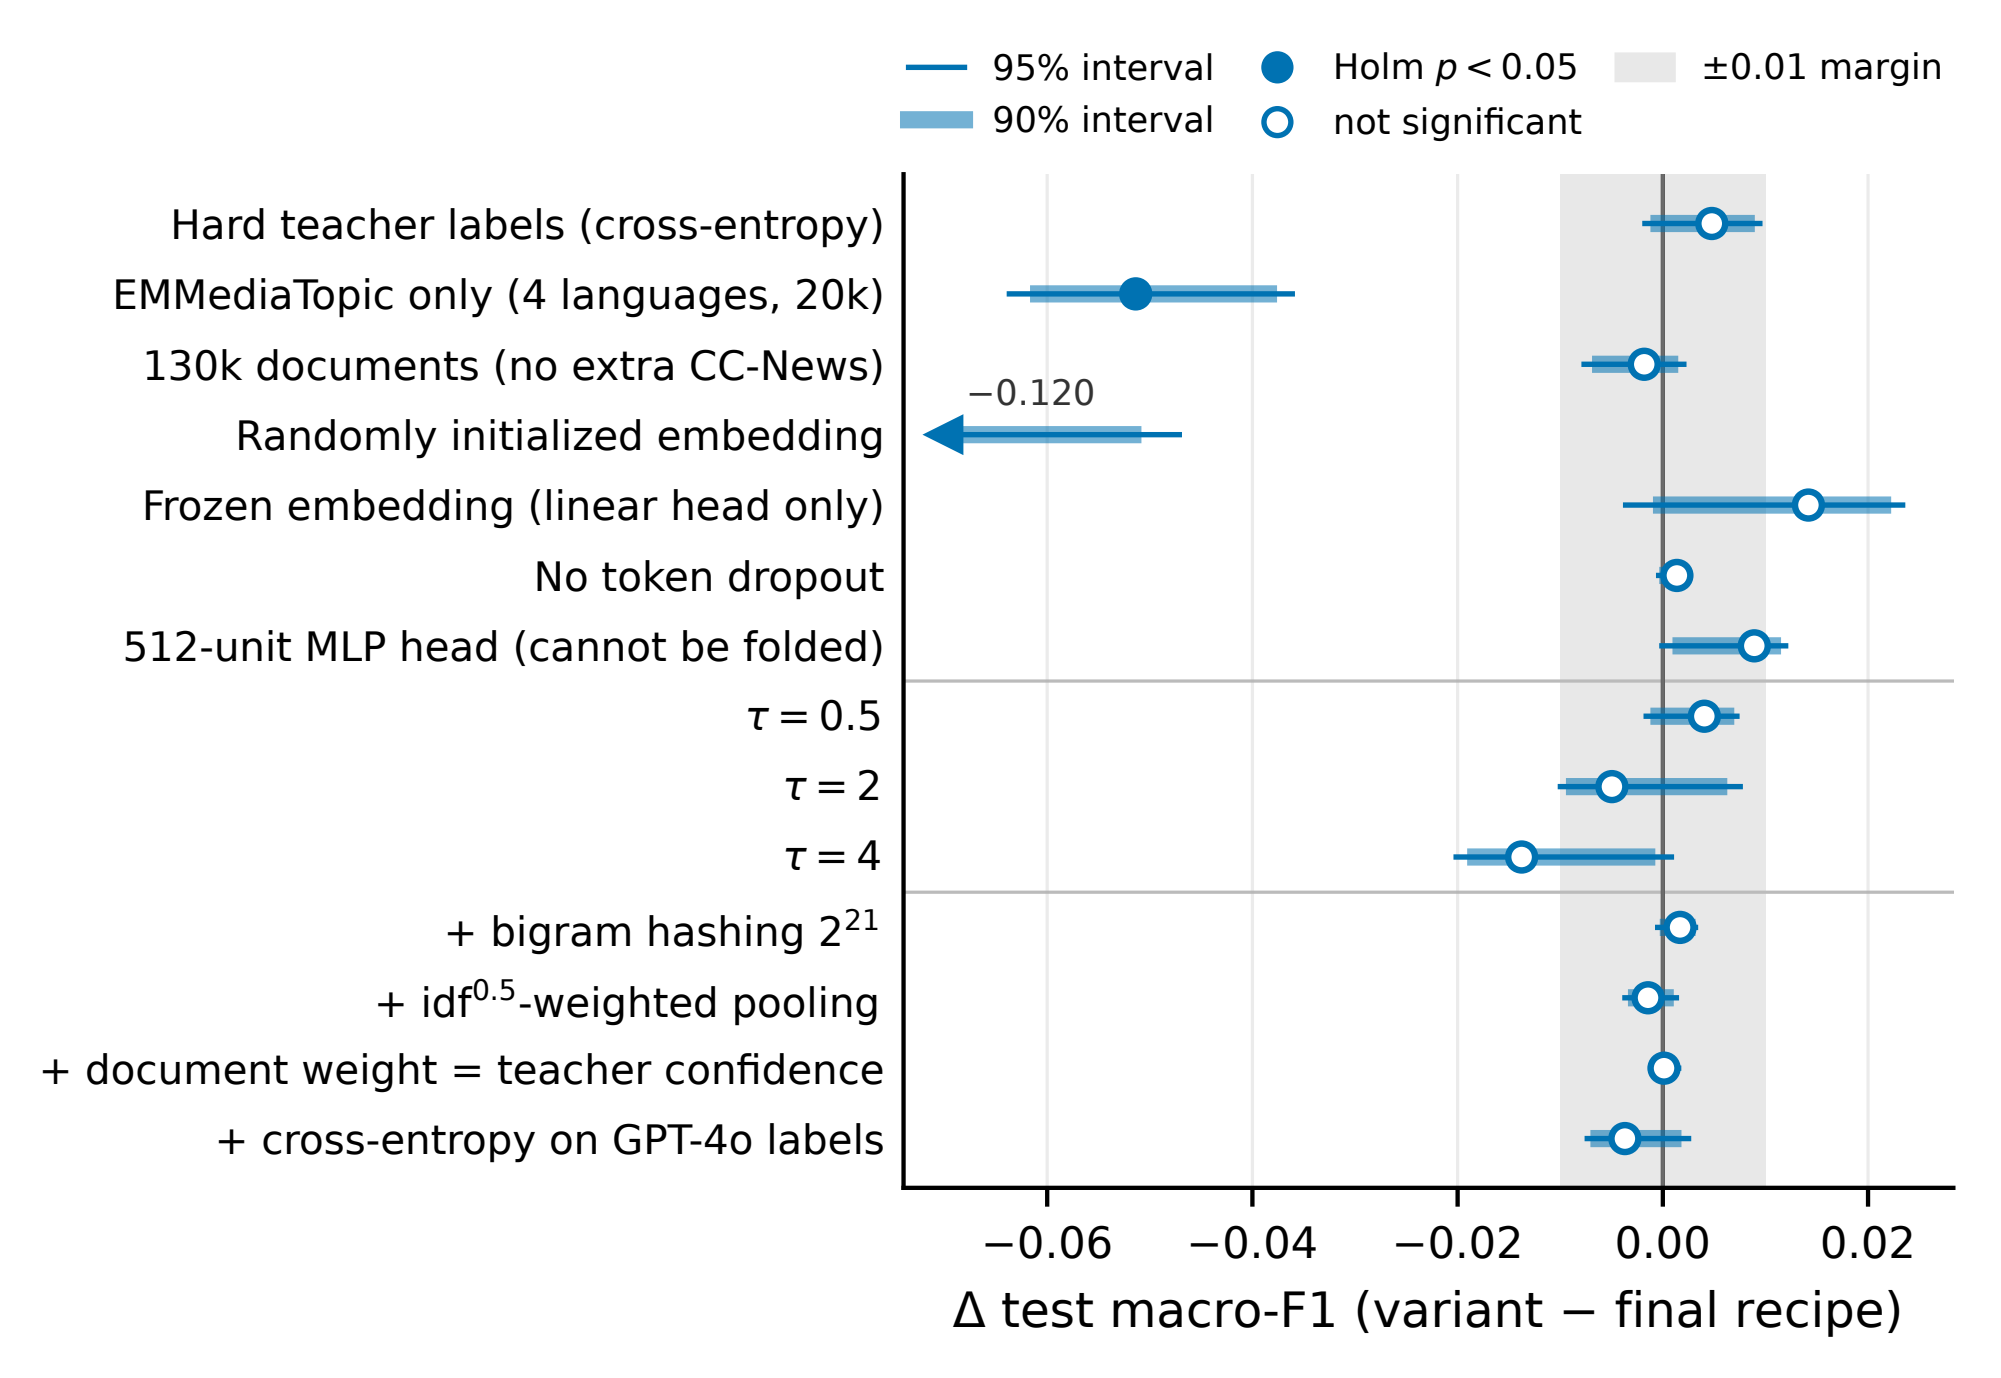
95% interval
90% interval
Holm $p < 0.05$
not significant
±0.01 margin
Hard teacher labels (cross-entropy)
EMMediaTopic only (4 languages, 20k)
130k documents (no extra CC-News)
Randomly initialized embedding
Frozen embedding (linear head only)
No token dropout
512-unit MLP head (cannot be folded)
$\tau = 0.5$
$\tau = 2$
$\tau = 4$
+ bigram hashing 2^{21}
+ idf$^{0.5}$-weighted pooling
+ document weight = teacher confidence
+ cross-entropy on GPT-4o labels
−0.120
−0.06
−0.04
−0.02
0.00
0.02
Δ test macro-F1 (variant − final recipe)